

COMPUTING OPTIMAL RANDOMIZED RESOURCE ALLOCATIONS FOR MASSIVE SECURITY GAMES

COURSE: EECE 580A – CYBER PHYSICAL SYSTEMS SECURITY

PRESENTER: JOSEPH ST. CYR

DATE: DECEMBER 2ND, 2019

SECURITY CONUNDRUM

- Providing security for transportation systems, computer networks, and other critical infrastructure is a large and growing problem.
- Central to many of these security problems is a resource allocation task.
 - E.g., A police force may have limited personnel to conduct patrols, operate checkpoints, conduct random searches, etc.
- *How can we efficiently allocate available resources to protect against a wide variety of potential threats?*

RESOURCE ALLOCATION RANDOMIZATION

- Adversarial aspect of security domains poses unique challenges for resource allocation.
- A motivated attacker can gather information about security measures by surveillance to plan more effective attacks.
- Predictable resource allocations may be exploited by an attacker, greatly reducing resource effectiveness.
- A better approach for deploying security resources is to use randomization to increase the uncertainty of potential attackers.

“COMBINATORIAL EXPLOSION”

- Game theory offers a more sophisticated approach to randomization.
- Allows the analyst to factor differential risks and costs into the game model.
- Incorporates predictions of how the attacker will respond to a given security policy.
- A significant limitation of existing solution methods:
 - Handle multiple resources by exhaustively enumerating all possible combinations (Normal Form).
 - Grows combinatorially in the number of resources, which makes it computationally infeasible to solve large problems.
 - 100 targets and 10 resources yields approx. 1.7×10^{13} allocation combinations (17 Trillion!!!)

GAME-THEORETIC MODELING

- Model the security resource allocation problem as a Stackelberg game.
 - Playing order to game (i.e., leader and follower).
 - Author's security games, the defender is a Stackelberg leader, and the attacker is a follower.
 - This models the capability of malicious attackers to employ surveillance in planning attacks.
- Defender Θ , and an attacker Ψ
- Each player has a set of possible pure strategies, denoted $\sigma_\Theta \in \Sigma_\Theta$ and $\sigma_\Psi \in \Sigma_\Psi$.
- A mixed strategy allows a player to play a probability distribution over pure strategies, denoted $\delta_\Theta \in \Delta_\Theta$ and $\delta_\Psi \in \Delta_\Psi$.

GAME-THEORETIC MODELING

- Payoffs for each player are defined over all possible joint pure strategy outcomes: $\Omega_\Theta : \Sigma_\Psi \times \Sigma_\Theta \rightarrow \mathbb{R}$ for the defender and similarly for the attacker.
- The payoff functions are extended to mixed strategies.
- Formally, the attacker's strategy in a Stackelberg security game is a function that selects a strategy in response to each leader strategy: $F_\Psi : \Delta_\Theta \rightarrow \Delta_\Psi$.

STRONG STACKELBERG EQUILIBRIUM

- Stackelberg equilibrium is a refinement of Nash equilibrium.
- Form of subgame perfect equilibrium, such that each player chooses a best-response.
- Eliminates Nash equilibrium profiles supported by non-credible threats off the equilibrium path.
- *DEFINITION 1.* A pair of strategies (δ_Θ, F_Ψ) form a Strong Stackelberg Equilibrium (SSE) if they satisfy the following:
 - The leader plays a best-response: $\Omega_\Theta(\delta_\Theta, F_\Psi(\delta_\Theta)) \geq \Omega_\Theta(\delta'_\Theta, F_\Psi(\delta'_\Theta)) \forall \delta'_\Theta \in \Delta_\Theta$.
 - The follower plays a best-response: $\Omega_\Psi(\delta_\Theta, F_\Psi(\delta_\Theta)) \geq \Omega_\Psi(\delta_\Theta, \delta_\Psi) \forall \delta_\Theta \in \Delta_\Theta, \delta_\Psi \in \Delta_\Psi$.
 - The follower breaks ties optimally: $\Omega_\Theta(\delta_\Theta, F_\Psi(\delta_\Theta)) \geq \Omega_\Theta(\delta_\Theta, \delta_\Psi) \forall \delta_\Theta \in \Delta_\Theta, \delta_\Psi \in \Delta^*_\Psi(\delta_\Theta)$, where $\Delta^*_\Psi(\delta_\Theta)$ is the set of follower best-responses, as above.

A COMPACT REPRESENTATION OF RESOURCES

- Let $T = \{t_1, \dots, t_n\}$ be a set of targets that may be attacked, corresponding to pure strategies for the attacker.
- The defender has a set of resources available to “cover” these targets, $R = \{r_1, \dots, r_m\}$.
- Associated with each target are four payoffs defining the possible outcomes for an attack on the target, as shown in Table 1.

Table 1: Example payoffs for an attack on a target.

	Covered	Uncovered
Defender	5	-20
Attacker	-10	30

A COMPACT REPRESENTATION OF RESOURCES

- There are two cases, depending on whether or not the target is covered by the defender.
 - The defender's payoff: Uncovered attack = $U_{\Theta}^u(t)$, Covered attack = $U_{\Theta}^c(t)$.
 - The attacker's payoff: Uncovered attack = $U_{\Psi}^u(t)$, Covered attack = $U_{\Psi}^c(t)$.
- A crucial feature:
 - Payoffs only depend on the identity of the attacked target and whether or not it is covered by the defender.
 - Does not matter whether or not any un-attacked target is covered or not.

A COMPACT REPRESENTATION OF RESOURCES

- Exploited by summarizing the payoff-relevant aspects:
 - Defender: coverage vector, C , that gives the probability that each target is covered, c_t .
 - Attacker: attack vector, A , that gives the probability that each target is attacked, a_t .

$$U_{\Theta}(C, A) = \sum_{t \in T} a_t \cdot (c_t \cdot U_{\Theta}^c(t) + (1 - c_t)U_{\Theta}^u(t)) \quad (1)$$

$$U_{\Theta}(t, C) = c_t U_{\Theta}^c(t) + (1 - c_t)U_{\Theta}^u(t) \quad (2)$$

$$\Gamma(C) = \{t : U_{\Psi}(t, C) \geq U_{\Psi}(t', C) \forall t' \in T\}. \quad (3)$$

$$\begin{aligned} \delta_{\Theta}^{1,2} + \delta_{\Theta}^{1,3} + \delta_{\Theta}^{1,4} &= c_1 \\ \delta_{\Theta}^{1,2} + \delta_{\Theta}^{2,3} + \delta_{\Theta}^{2,4} &= c_2 \\ \delta_{\Theta}^{1,3} + \delta_{\Theta}^{2,3} + \delta_{\Theta}^{3,4} &= c_3 \\ \delta_{\Theta}^{1,4} + \delta_{\Theta}^{2,4} + \delta_{\Theta}^{3,4} &= c_4 \end{aligned} \quad (4)$$

- In a strong Stackelberg equilibrium:
 - Attacker selects the target in the attack set with maximum payoff for the defender.
 - Let t^* denote optimal target.
 - The expected SSE payoff for the defender is $\hat{U}_{\Theta}(C) = U_{\Theta}(t^*, C)$, and $\hat{U}_{\Psi}(C) = U_{\Psi}(t^*, C)$ for the attacker.

“ERASER” SOLUTION ALGORITHM

- The **ERASER** algorithm (**E**fficient **R**andomized **A**llocation of **S**Ecurity **R**esources) takes as input a security game in compact form and solves for an optimal coverage vector corresponding to a SSE strategy for the defender.
- The algorithm is a mixed-integer linear program (MILP).
- C and A are mutual best-responses in any optimal solution.

$$\max_{a_t \in \{0, 1\}} d \quad (5)$$

$$\sum_{t \in T} a_t = 1 \quad (6)$$

$$c_t \in [0, 1] \quad \forall t \in T \quad (7)$$

$$\sum_{t \in T} c_t \leq m \quad (8)$$

$$d - U_\Theta(t, C) \leq (1 - a_t) \cdot Z \quad \forall t \in T \quad (9)$$

$$0 \leq k - U_\Psi(t, C) \leq (1 - a_t) \cdot Z \quad \forall t \in T \quad (10)$$

$$(11)$$

EXPLOITING PAYOFF STRUCTURE

- Consider a class of security games in which the defender always benefits by having additional resources covering an attacked target, while the attacker is always worse off attacking a more heavily defended target.
- Formally, we restrict payoff functions so that:
 - $U^u_\Theta(t) < U^c_\Theta(t)$ and $U^u_\Psi(t) > U^c_\Psi(t)$ for all t .

OBSERVATION 1. *All else equal, increasing c_t for any target not in $\Gamma(C)$ has no effect on $\hat{U}_\Theta(C)$ or $\hat{U}_\Psi(C)$.*

OBSERVATION 2. *If $\Gamma(C) \subset \Gamma(C')$ and $c_t = c'_t$ for all $t \in \Gamma(C)$ then $\hat{U}_\Theta(C) \leq \hat{U}_\Theta(C')$.*

OBSERVATION 3. *If $\hat{U}_\Psi(C) = x$, then $c_t \geq \frac{x - U^u_\Psi(t)}{U^c_\Psi(t) - U^u_\Psi(t)}$ for every target t with $U^u_\Psi(t) > x$.*

“ORIGAMI” SOLUTION ALGORITHM

- Observations 1-3 are exploited in the **ORIGAMI** algorithm (**O**ptimizing **R**esources **I**n **G**Ames using **M**aximal **I**ndifference).
 - Directly compute the attack set for the attacker, using the indifference equation in Observation 3.
 - Starting with a target that has maximal $U_{\Psi}^u(t)$, the attack set is expanded at each iteration in order of decreasing $U_{\Psi}^u(t)$.
 - Each time the attack set is expanded, the coverage of each target is updated to maintain the indifference of attacker payoffs within the attack set.

“ORIGAMI” SOLUTION ALGORITHM

- Two termination conditions:
 - 1 - Occurs when adding the next target to the attack set requires more total coverage probability than the defender has resources available.
 - 2 - Occurs when any target t is covered with probability 1.

Algorithm 1 ORIGAMI

```

targets  $\leftarrow T$  sorted by  $U_{\Psi}^u(t)$ 
payoff[t]  $\leftarrow U_{\Psi}^u(t)$ , coverage[t]  $\leftarrow 0$ 
left  $\leftarrow m$ , next  $\leftarrow 2$ 
covBound  $\leftarrow -\infty$ 
while next  $\leq n$  do
    addedCov[t]  $\leftarrow \frac{\text{payoff}[next] - U_{\Psi}^u(t)}{U_{\Psi}^c(t) - U_{\Psi}^u(t)}$  - coverage[t]
    if coverage[t] + addedCov[t]  $\geq 1$  then
        covBound  $\leftarrow \text{Max}(\text{covBound}, U_{\Psi}^c(t))$ 
    end if
    if covBound  $\geq -\infty$  OR  $\sum_{t \in T} \text{addedCov}[t] \leq \text{left}$  then
        BREAK
    end if
    coverage[t]  $\leftarrow \text{coverage}[t] + \text{addedCov}[t]$ 
    left  $\leftarrow \sum_{t \in T} \text{addedCov}[t]$ 
    next++
end while
ratio[t]  $\leftarrow \frac{1}{U_{\Psi}^u(t) - U_{\Psi}^c(t)}$ 
coverage[t]  $\leftarrow \frac{\text{ratio}[t] \cdot \text{left}}{\sum_{t \in T} \text{ratio}[t]}$ 
if coverage[t]  $\geq 1$  then
    covBound  $\leftarrow \text{Max}(\text{covBound}, U_{\Psi}^c(t))$ 
end if
if covBound  $\geq -\infty$  then
    coverage[t]  $\leftarrow \frac{\text{covBound} - U_{\Psi}^u(t)}{U_{\Psi}^c(t) - U_{\Psi}^u(t)}$ 
end if

```

OTHER ALGORITHMS

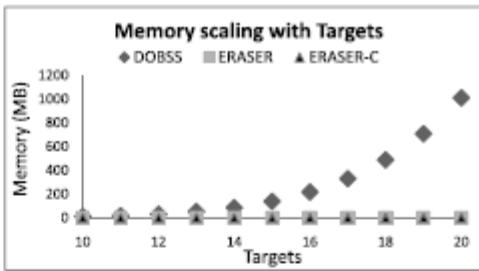
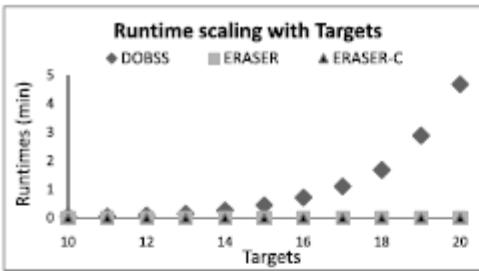
- ERASER-C

- An extension of ERASER which adds the capability to represent certain kinds of resource and scheduling constraints, motivated by the real example domains described previously.

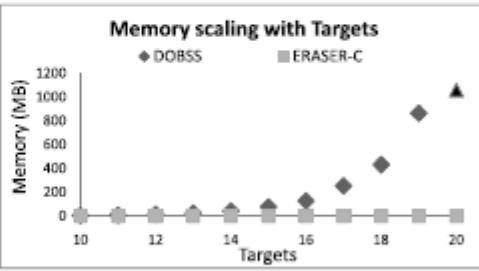
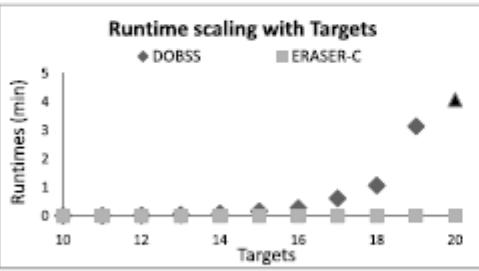
- ORIGAMI-MILP

- Similar to ERASER, but does not optimize the defender's payoff. Instead, it minimizes the attacker's payoff, and adds a constraint that restricts c_t for any t not in $\Gamma(C)$ to 0, consistent with Observation 1. This constraint forces the attack set to include the maximal number of targets.

EXPERIMENTAL RESULTS

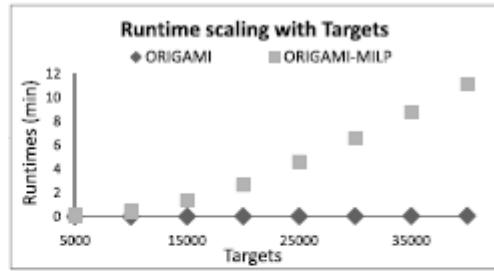
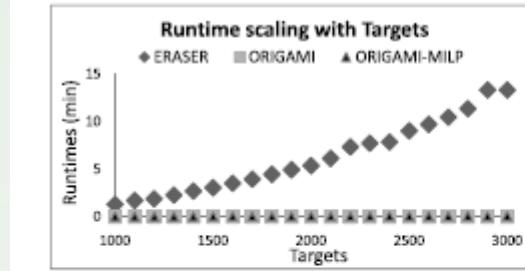


(a) Runtimes for DOBSS, (b) Memory use of DOBSS,
ERASER, and ERASER-C (c) Runtimes for DOBSS and
ERASER-C (d) Memory use of DOBSS and
ERASER-C



(a) Runtimes for DOBSS, (b) Memory use of DOBSS,
ERASER, and ERASER-C (c) Runtimes for DOBSS and
ERASER-C (d) Memory use of DOBSS and
ERASER-C

Figure 1: Runtime and memory scaling



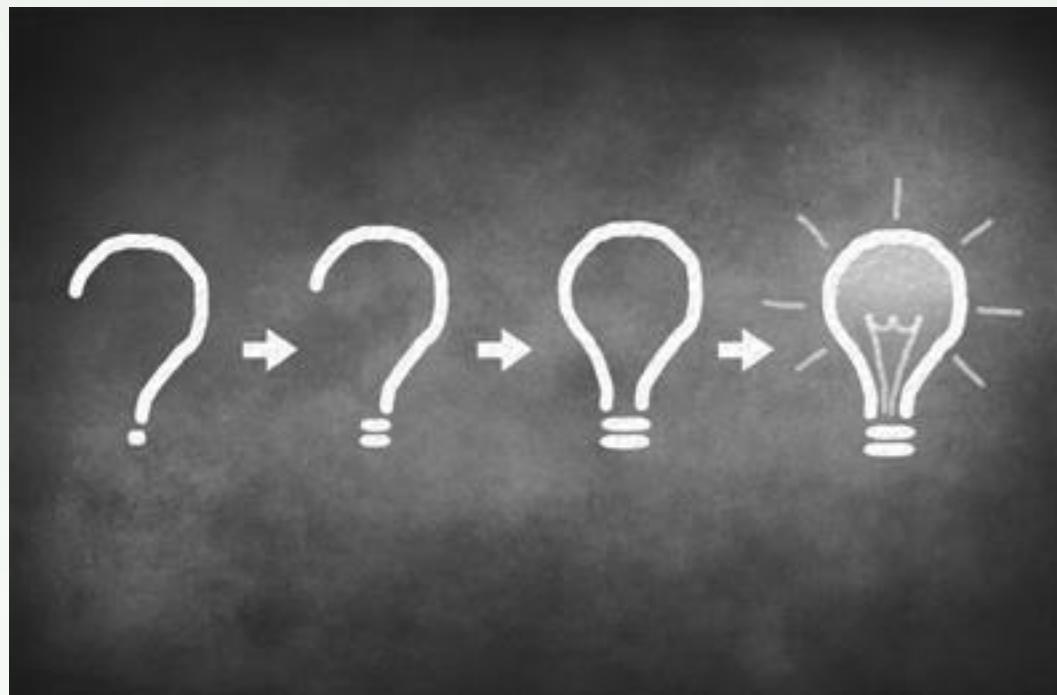
(a) (b)

Figure 2: Runtime scaling of ERASER, ORIGAMI, and ORIGAMI-MILP

	Actions	DOBSS	ERASER (-C)
LAX (6 canines)	784	0.94s	0.23s
FAMS (small)	~6,000	4.74s	0.09s
FAMS (large)	~85,000	435.6s*	1.57s

Table 2: Runtimes on real data.

QUESTIONS & DISCUSSION



REFERENCES

1. Kiekintveld, Christopher & Jain, Manish & Tsai, Jason & Pita, James & Ordóñez, Fernando & Tambe, Milind. (2009). Computing optimal randomized resource allocations for massive security games. *AMAAS-2009 Conf.* 2. 689-696. [10.1145/1558013.1558108](https://doi.org/10.1145/1558013.1558108).